

Module: XML et les bases de données



An Open Source Native XML Database

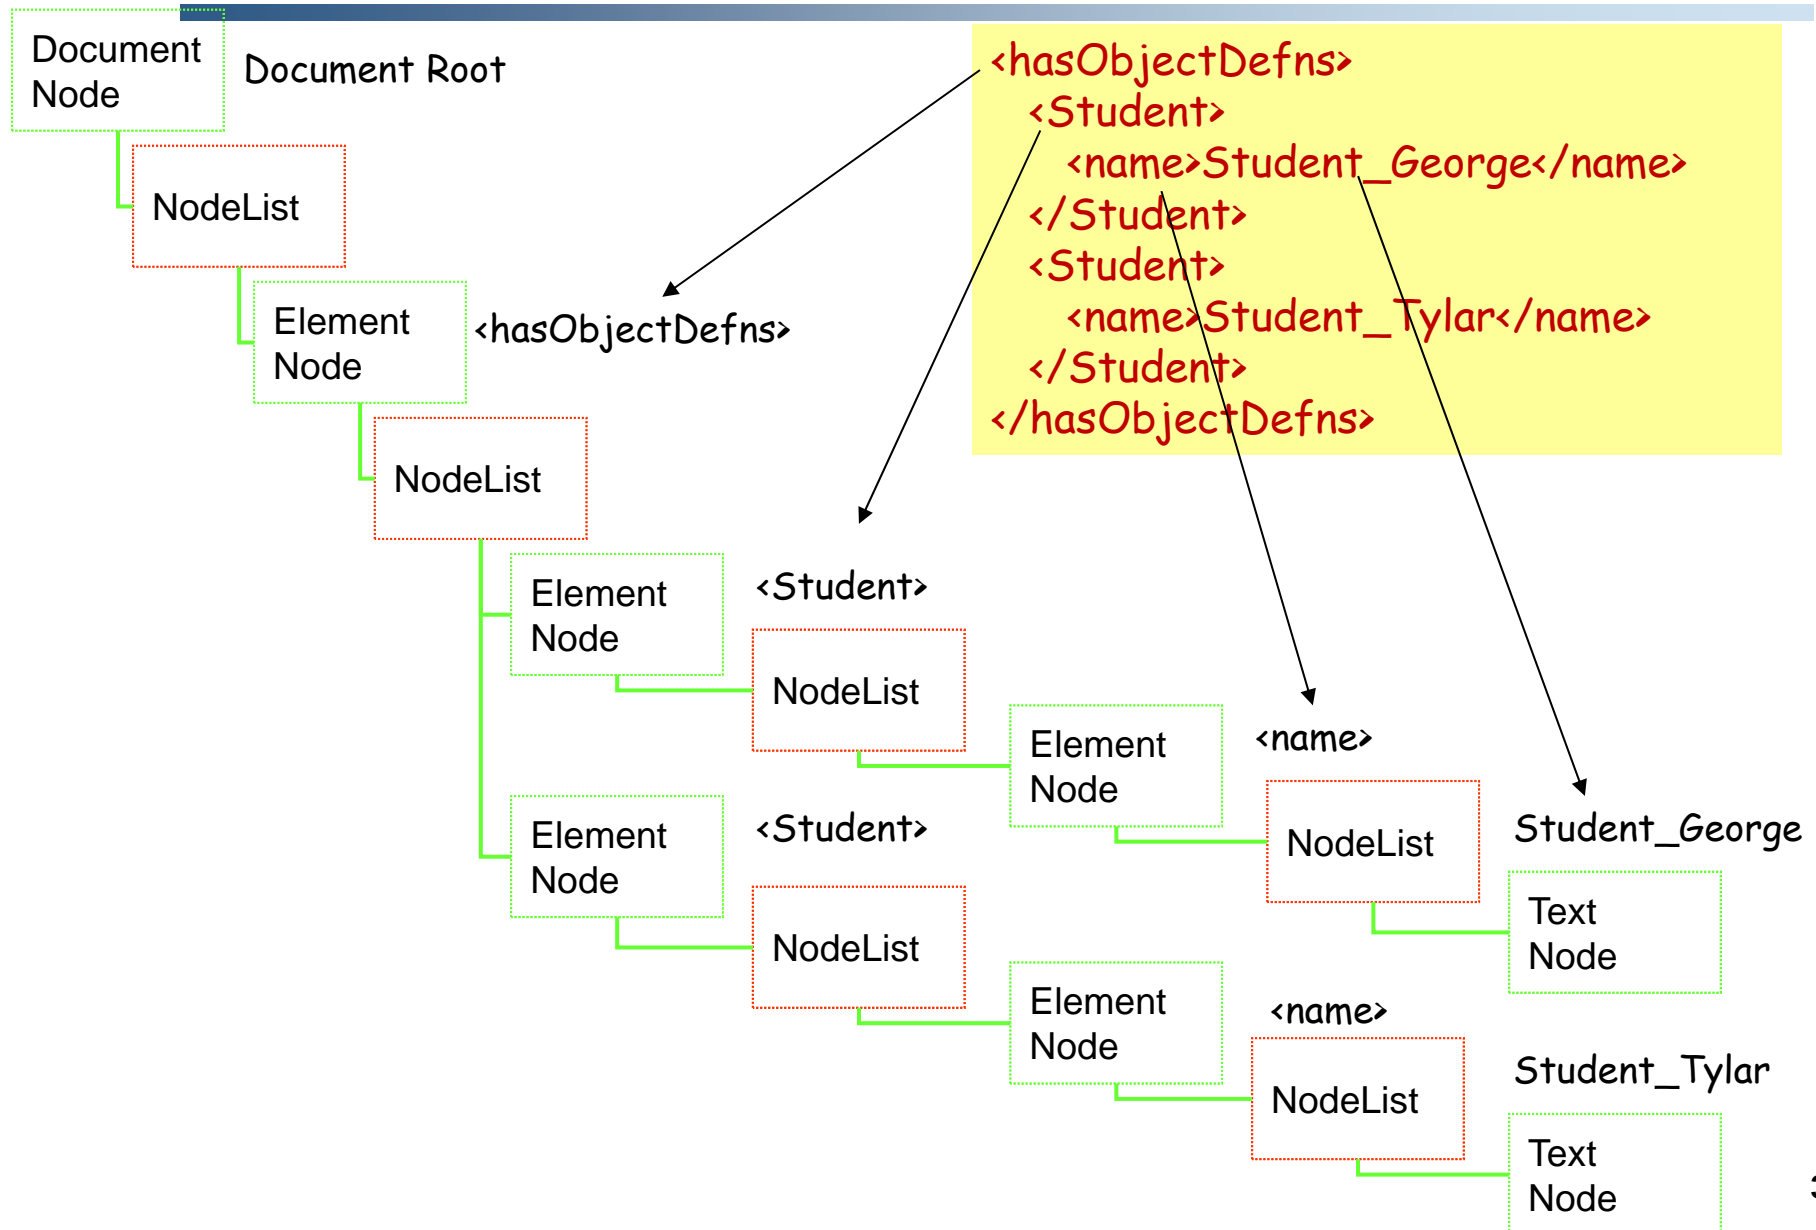
Houda Chabbi Drissi

houda.chabbi@hefr.ch

eXist

- eXist an open source native XML database, written in Java (platform independent), with an XQuery query processor.
- First version: 2000
- It stores native XML data in:
 - B+-trees and paged files,
 - Document nodes in persistent DOM trees.
- Documents can be divided into collections. The collections can be arranged in a collection hierarchy (similar to a regular file system).

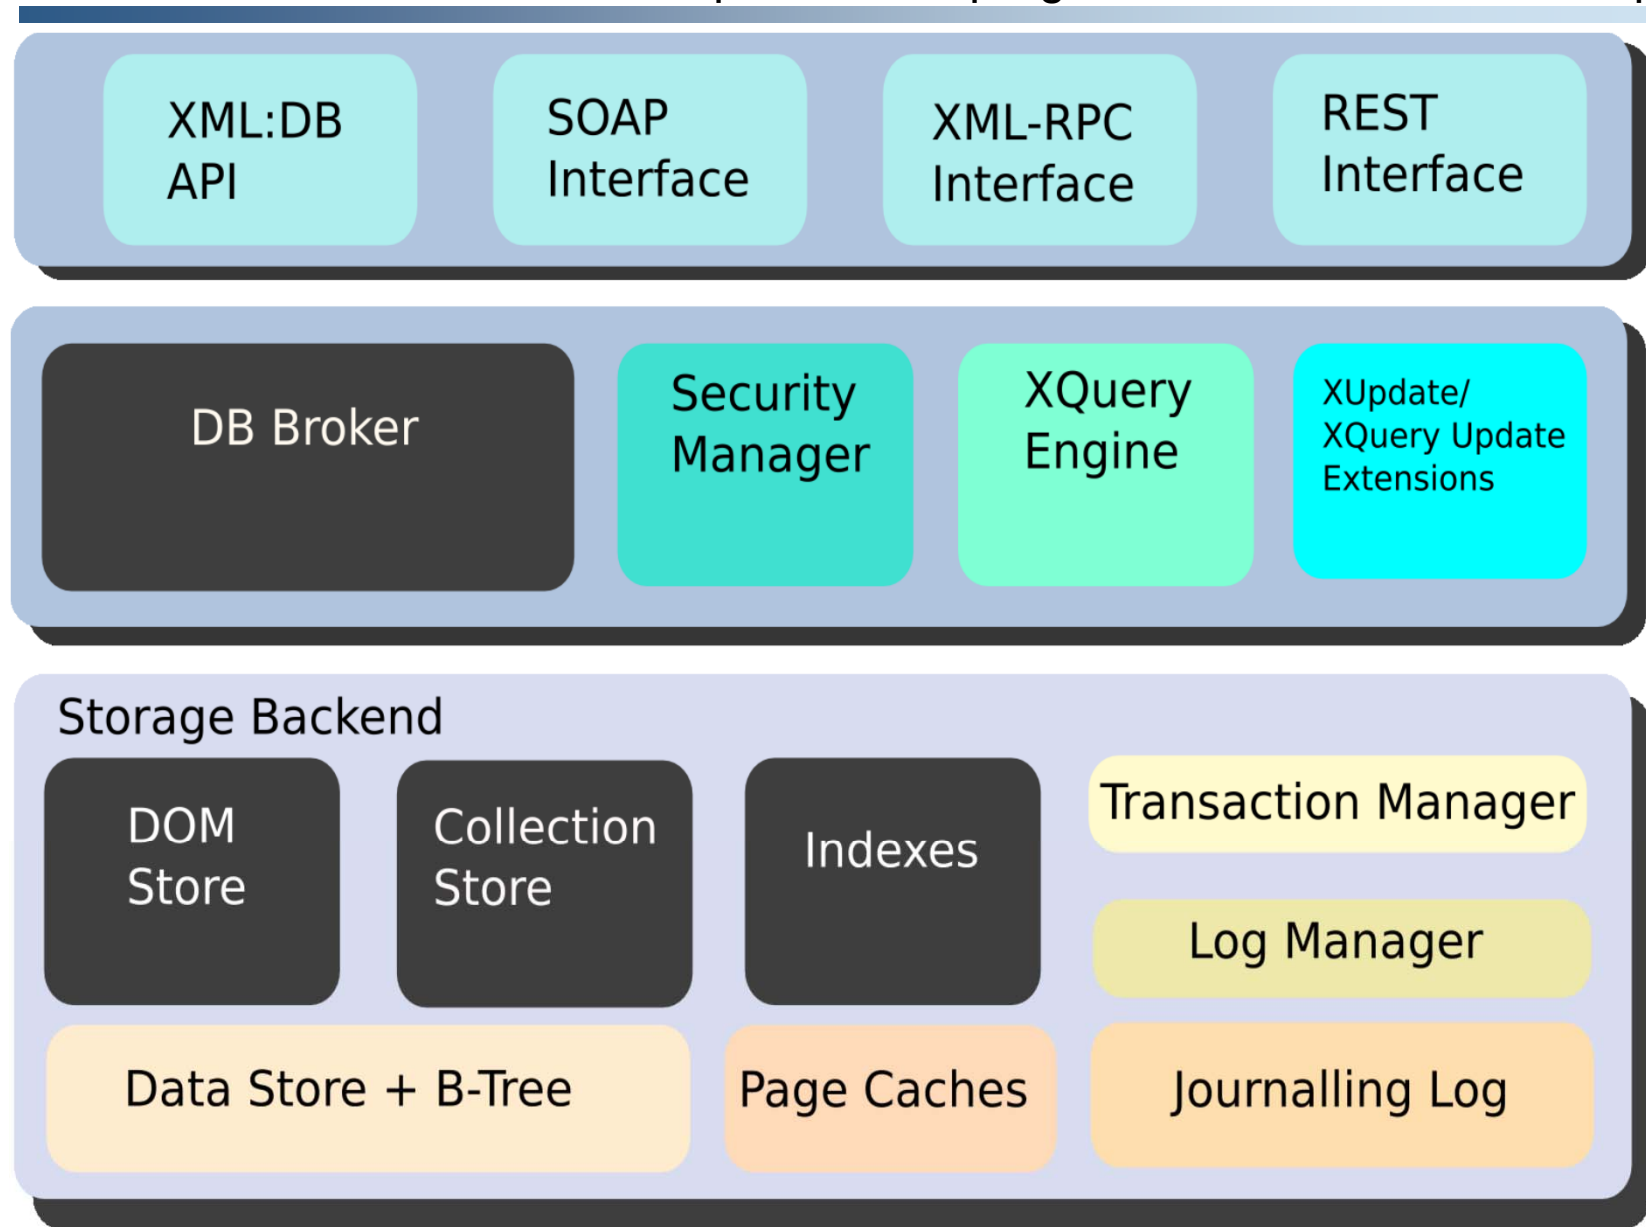
Document Object Model: DOM



eXist

- Transaction support: No
- Authorization: Unix like, permissions at collection and document level
- XML Standards that are supported:
 - ✓ XPath
 - ✓ Xquery
 - ✓ XUpdate
- Comes with great client GUI interface
- Types of indexes: Structural, Fulltext

Source: <http://www.xmlprague.cz/2006/slides06/meier.pdf>



Index and Data Organization

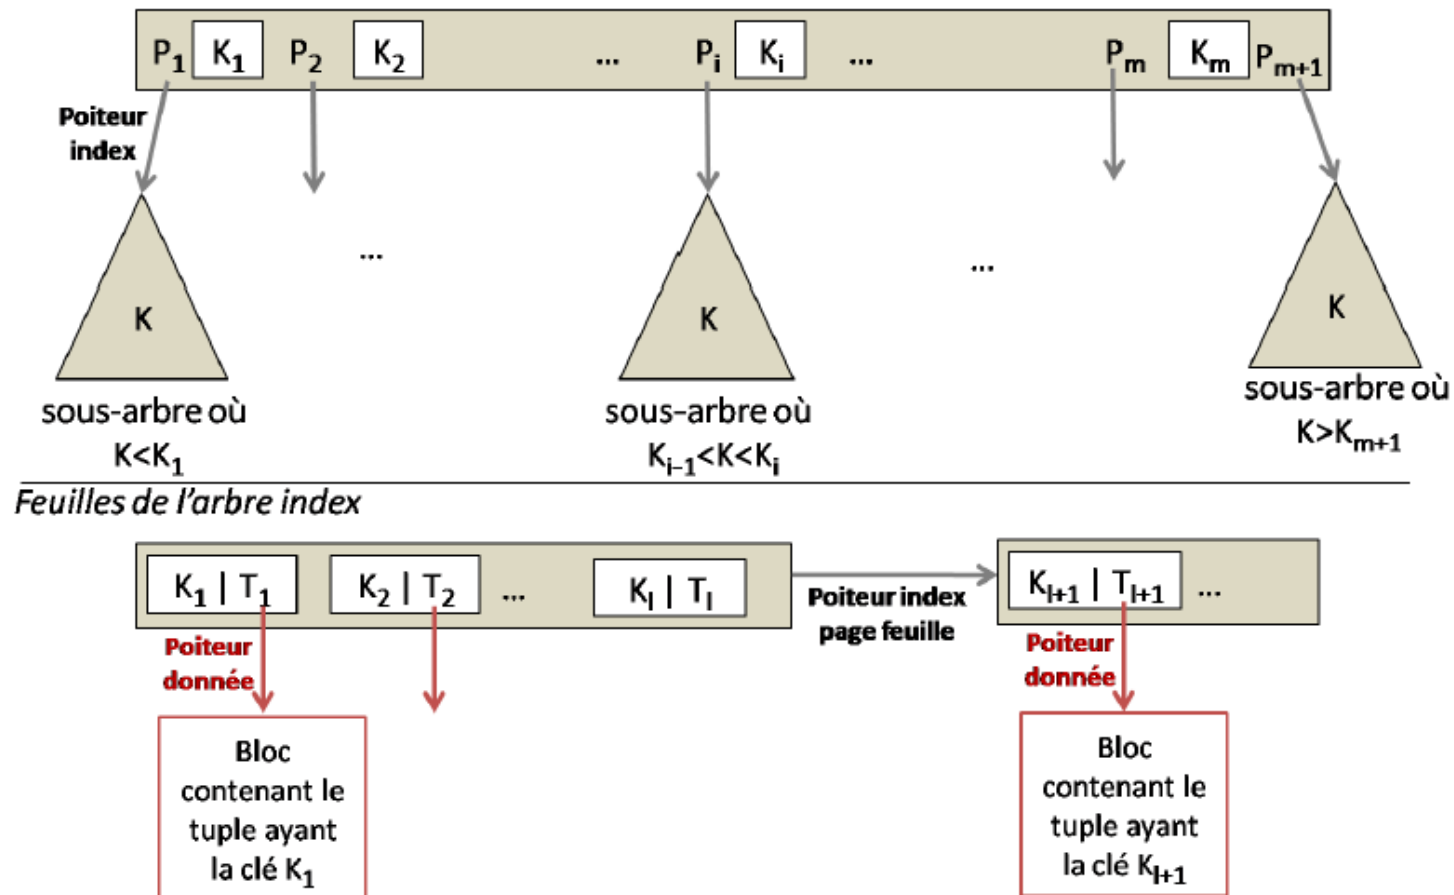
- eXist uses four index files at the XML storage backend:
 - **collections.dbx** - manages the collection hierarchy
 - **dom.dbx** - collects nodes in a paged file and associates unique node identifiers to the actual nodes
 - **elements.dbx** - indexes elements and attributes
 - **words.dbx** - keeps track of word occurrences and is used by the fulltext search extensions → old version replaced by a new one based on **lucene**.

All based on **B⁺-trees**

B+ arbre

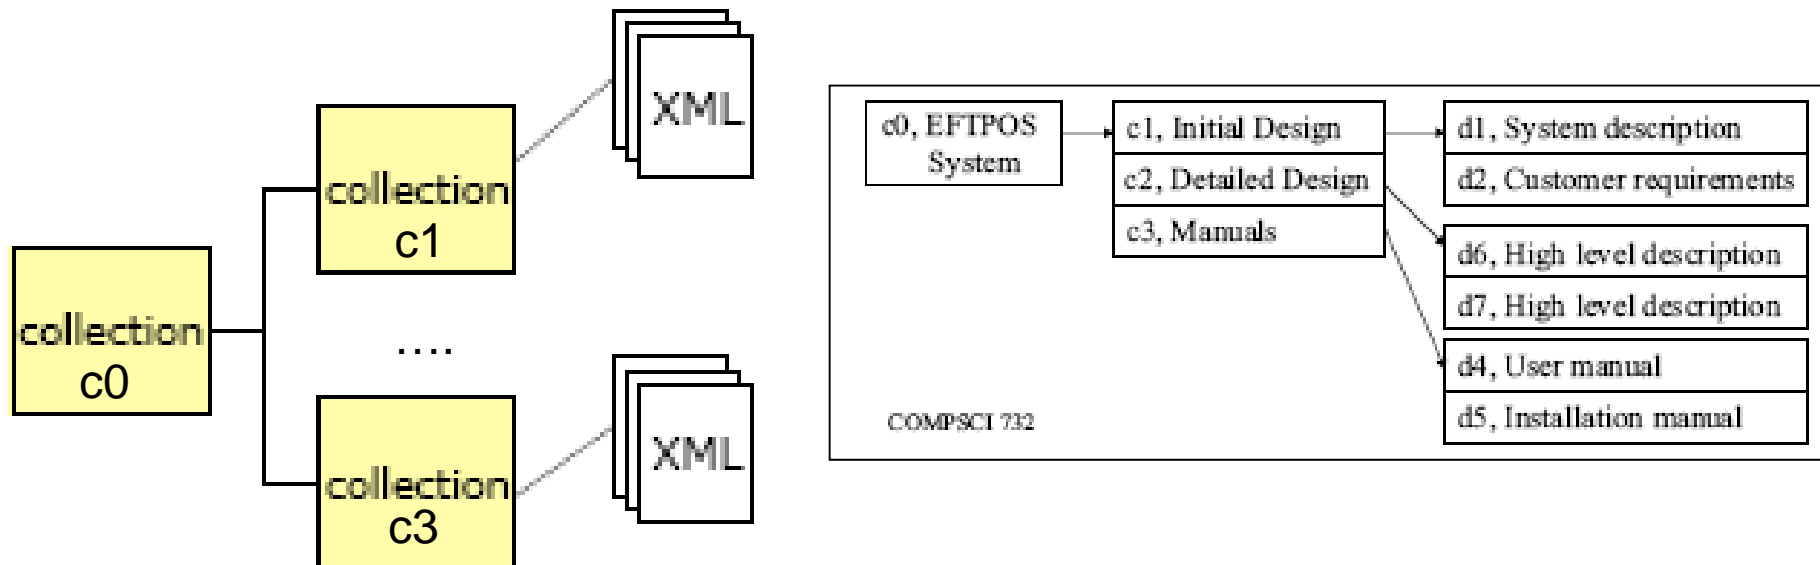
■ Un B+-arbre d'ordre d = un B-arbre où

- Les nœuds internes ne contiennent que les clés et les pointeurs d'index
- Les feuilles de l'arbre contiennent toutes les clés et les pointeurs de données



collection.dbx

- Manages the collection hierarchy and maps collection names to collection objects.



Indexes for **elements, attributes and keywords** are organized by collection and not by document

dom.dbx

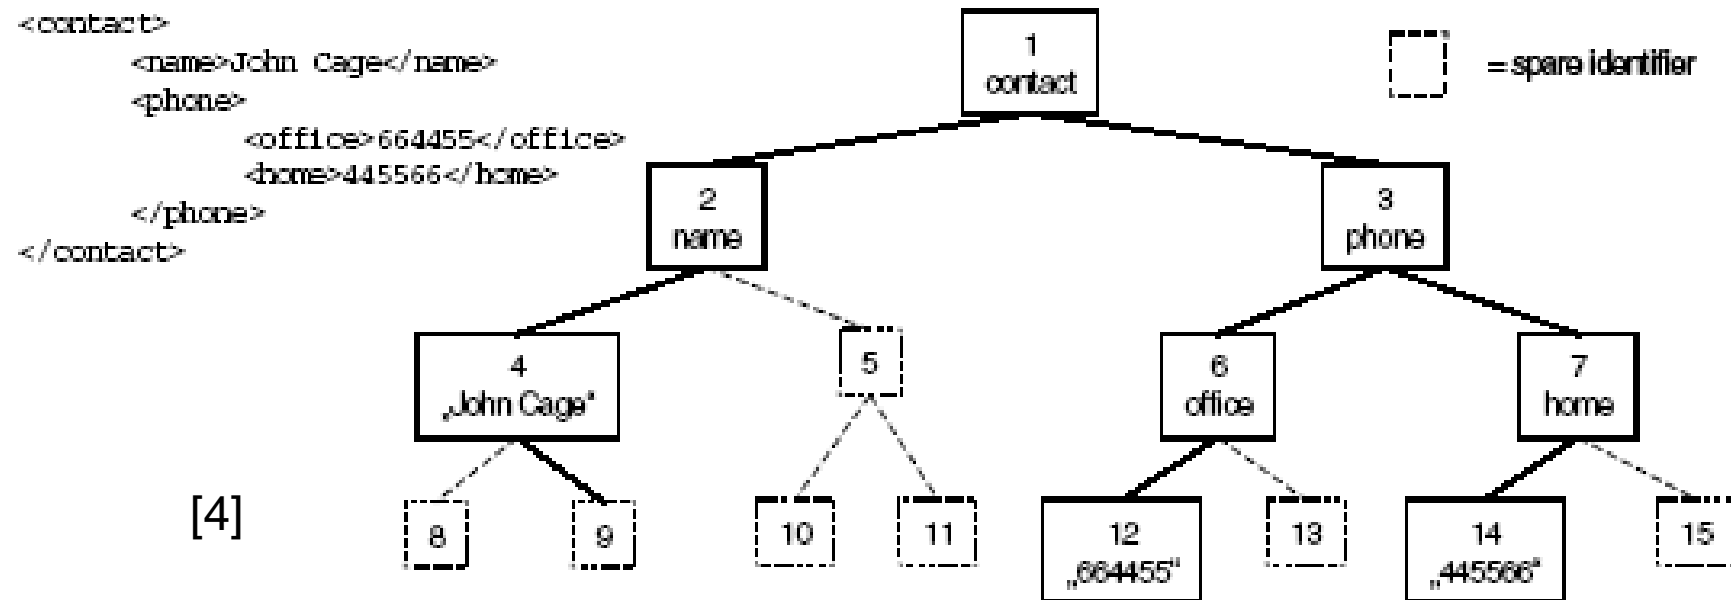
- All document nodes are stored according to the W3C's DOM (Document Object Model).
- In eXist DOM nodes are associated to unique node identifiers.
- Associate the unique node identifiers of top-level elements in a given document to the node's storage address in the data pages.

XML Index by Numbering

- A numbering scheme assigns a unique identifier to each node in an XML document
 - Old approach: Level-Order Numbering (or Virtual Node Numbering)
 - New approach: Dynamic Level Numbering (DLN) (2006)
- Certain relationships can be determined immediately

Level-Order Numbering

- XML document is modeled as a complete k-ary tree (each node-root has K children)
- Unique identifier assigned to each node by traversing tree document in level-order




Level-Order Numbering: benefits

- A node's parent identifier can be computed with

$$\text{parent}_i = \left\lfloor \frac{(i-2)}{k} + 1 \right\rfloor \quad [4]$$

- Exemple: $\text{parent}_7 = \left\lfloor \frac{(7-2)}{2} + 1 \right\rfloor = \left\lfloor 3.5 \right\rfloor = 3$



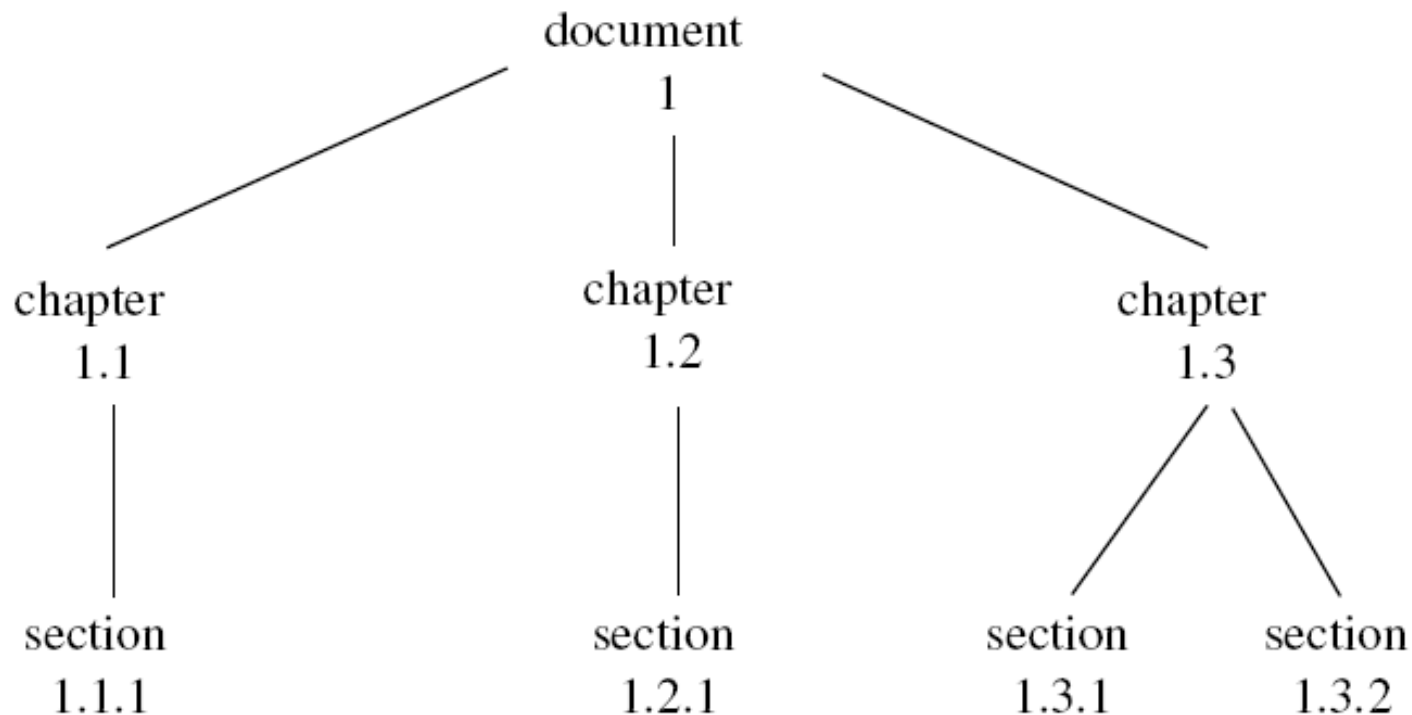
Level-Order Numbering: Disadvantage

- Typical documents have a **small number of top level** elements (e.g. chapters of a book), whereas **the deeper elements contain many more children** (e.g. paragraphs in a chapter)
 - Spare identifiers need to be inserted at the top level elements (to be complete)
- Not update friendly: node insertions may trigger a complete renumbering of the tree

Dynamic Level Numbering

- Path-based identification scheme
- Node IDs consist of the ID of the parent node as prefix and a level value: **1, 1.1, 1.2, 1.2.1 . . .**
- Between two nodes 1.1 and 1.2, a new node can be inserted as 1.1/1


Figure: Node IDs assigned by DLN



Dynamic Level Numbering

- To insert a node before **1.1**, we assign level-value 0, i.e. **1.0/1**
- The next node inserted after **1.0/1** then gets ID **1.0/0/1**

Figure: Node tree after updates



Dynamic Level Numbering: Disadvantage

- Depending on nesting depth, IDs can become very long
- Main challenge: find an efficient binary encoding

Example: IDs Picked From a Real TEI Document

```
1 <div6 exist:id="1.46.9.7.8" xml:id="JG10229">
2   <head exist:id="1.46.9.7.8.7">
3     <name exist:id="1.46.9.7.8.7.3" type="Dramenfigur">FAUST</name>
4     <stage exist:id="1.46.9.7.8.7.4">unruhig</stage>
5     <stage exist:id="1.46.9.7.8.7.6"> auf seinem Sessel am Pulten</stage>
6   </head>
7   <sp exist:id="1.46.9.7.8.8" xml:id="JG10230">
8     <lg exist:id="1.46.9.7.8.8.4">
9       <l exist:id="1.46.9.7.8.8.4.7" n="1" part="N">Hab nun ach die Philosophie</l>
10      <l exist:id="1.46.9.7.8.8.4.8" part="N">Medizin und Juristerey,</l>
11      <l exist:id="1.46.9.7.8.8.4.9" part="N">Und leider auch die Theologie</l>
12      <l exist:id="1.46.9.7.8.8.4.10" part="N">
13        Durchaus<note exist:id="1.46.9.7.8.8.4.10.4" place="unspecified" anchored="yes">
14          <hi exist:id="1.46.9.7.8.8.4.10.4.4">Durchaus]</hi> Vollständig.</note>
15          studirt mit heisser Müh.
16        </l>
17      </lg>
18    </sp>
19  </div6>
```


Annexe Annexe: Binary Encoding of a Level Value

Variable-length encoding using fixed-size units (currently: 4 bits)

Units	Bit pattern	ID range
1	0XXX	1..7
2	10XX XXXX	8..71
3	110X XXXX XXXX	72..583
4	1110 XXXX XXXX XXXX	584..4679
5	1111 0XXX XXXX XXXX XXXX	4680..37447
6	1111 10XX XXXX XXXX XXXX XXXX	37448..299591

Annexe: Binary Encoding of a Level Value

Variable-length encoding using fixed-size units (currently: 4 bits)

Units	Bit pattern	ID range
1	0XXX	1..7
2	10XX XXXX	8..71
3	110X XXXX XXXX	72..583
4	1110 XXXX XXXX XXXX	584..4679
5	1111 0XXX XXXX XXXX XXXX	4680..37447
6	1111 10XX XXXX XXXX XXXX XXXX	37448..299591

A DLN is encoded as a sequence of stream-encoded level values separated by a 0-bit.

ID	Bit string	Bits
1.3	0001 0 0011	9
1.80	0001 0 1100 0000 1001	17
1.10000.1	0001 0 11110001010011001001 0 0001	30

Annexe Annexe: Binary Encoding of a Level Value

Variable-length encoding using fixed-size units (currently: 4 bits)

Units	Bit pattern	ID range
1	0XXX	1..7
2	10XX XXXX	8..71
3	1 10X XXXX XXXX	72..583
4	1110 XXXX XXXX XXXX	584..4679
5	1111 0XXX XXXX XXXX XXXX	4680..37447
6	1111 10XX XXXX XXXX XXXX XXXX	37448..299591

A DLN is encoded as a sequence of stream-encoded level values separated by a 0-bit.

ID	Bit string	Bits
1.3	0001 0 0011	9
1.80	0001 0 1 100 0000 1001	17
1.10000.1	0001 0 11110001010011001001 0 0001	30

Annexe Annexe: Binary Encoding of a Level Value

Variable-length encoding using fixed-size units (currently: 4 bits)

Units	Bit pattern	ID range
1	0XXX	1..7
2	10XX XXXX	8..71
3	110X XXXX XXXX	72..583
4	1110 XXXX XXXX XXXX	584..4679
5	1111 0XXX XXXX XXXX XXXX	4680..37447
6	1111 10XX XXXX XXXX XXXX XXXX	37448..299591

A DLN is encoded as a sequence of stream-encoded level values separated by a 0-bit.

ID	Bit string	Bits
1.3	0001 0 0011	9
1.80	0001 0 1100 0000 1001	17
1.10000.1	0001 0 1111 0001 0100 1100 1001 0 0001	30

Annexe Annexe: Binary Encoding of a Level Value

Variable-length encoding using fixed-size units (currently: 4 bits)

Units	Bit pattern	ID range
1	0XXX	1..7
2	10XX XXXX	8..71
3	110X XXXX XXXX	72..583
4	1110 XXXX XXXX XXXX	584..4679
5	1111 0XXX XXXX XXXX XXXX	4680..37447
6	1111 10XX XXXX XXXX XXXX XXXX	37448..299591

A DLN is encoded as a sequence of stream-encoded level values separated by a 0-bit.

ID	Bit string	Bits
1.3	0001 0 0011	9
1.80	0001 0 1100 0000 1001	17
1.10000.1	0001 0 11110001010011001001 0 0001	30

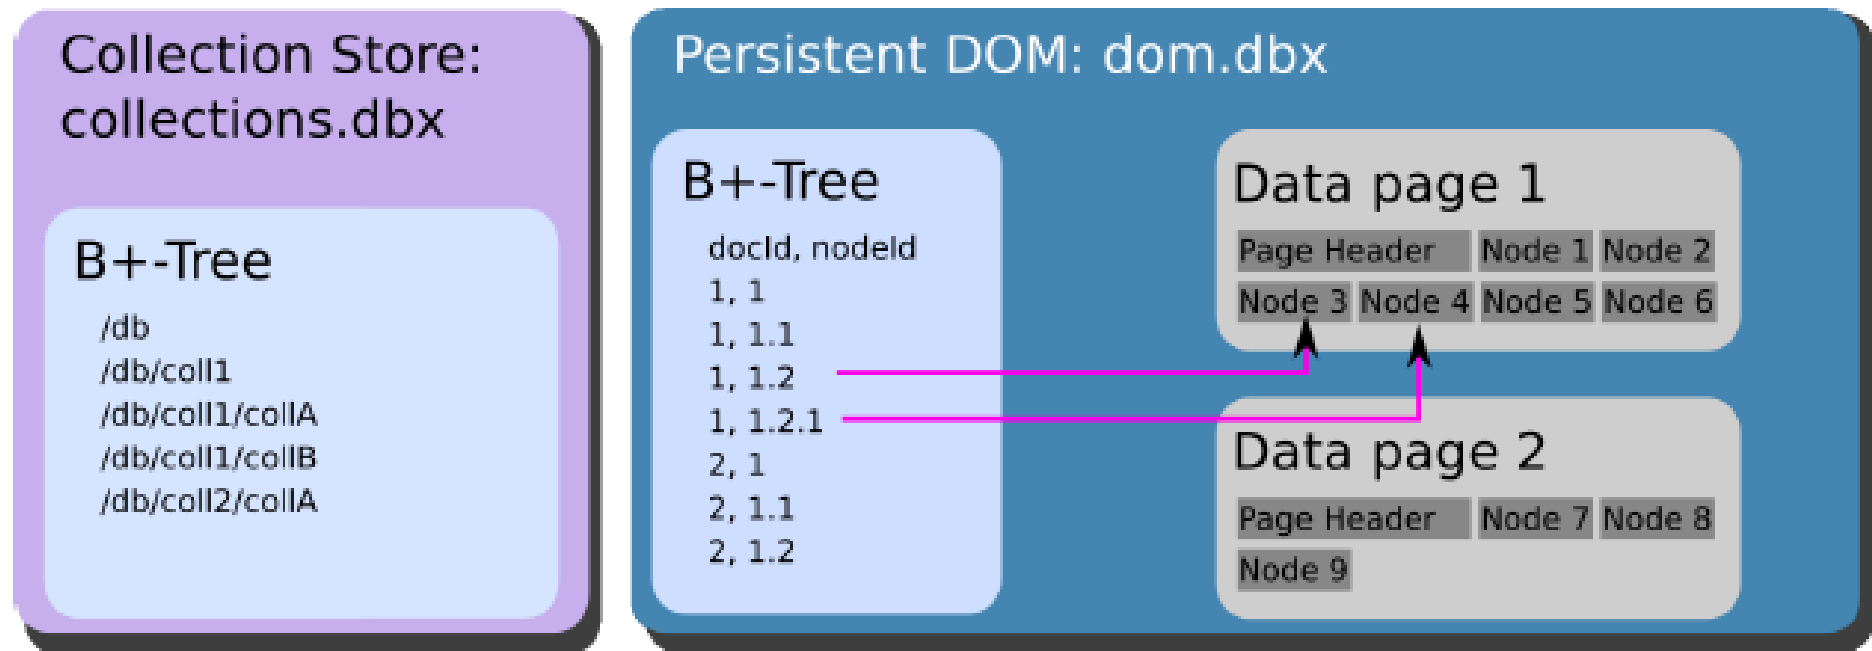
- To store a DLN with subvalue, we use a 1-bit to separate subvalues
- Level values are separated with a 0-bit as before
- Example: 1.1007 is encoded as 0001 0 0001 1 1000

Dynamic Level Numbering: notes

- Repeatedly insert a node in front of the first child of an element IDs grow very fast: 1.1/0/1, 1.1/0/0/1, 1.1/0/0/0/1
- To handle this edge case, eXist triggers a defragmentation run after several insertions
- Defragmentation is required here anyway to reduce the growth of dom.dbx

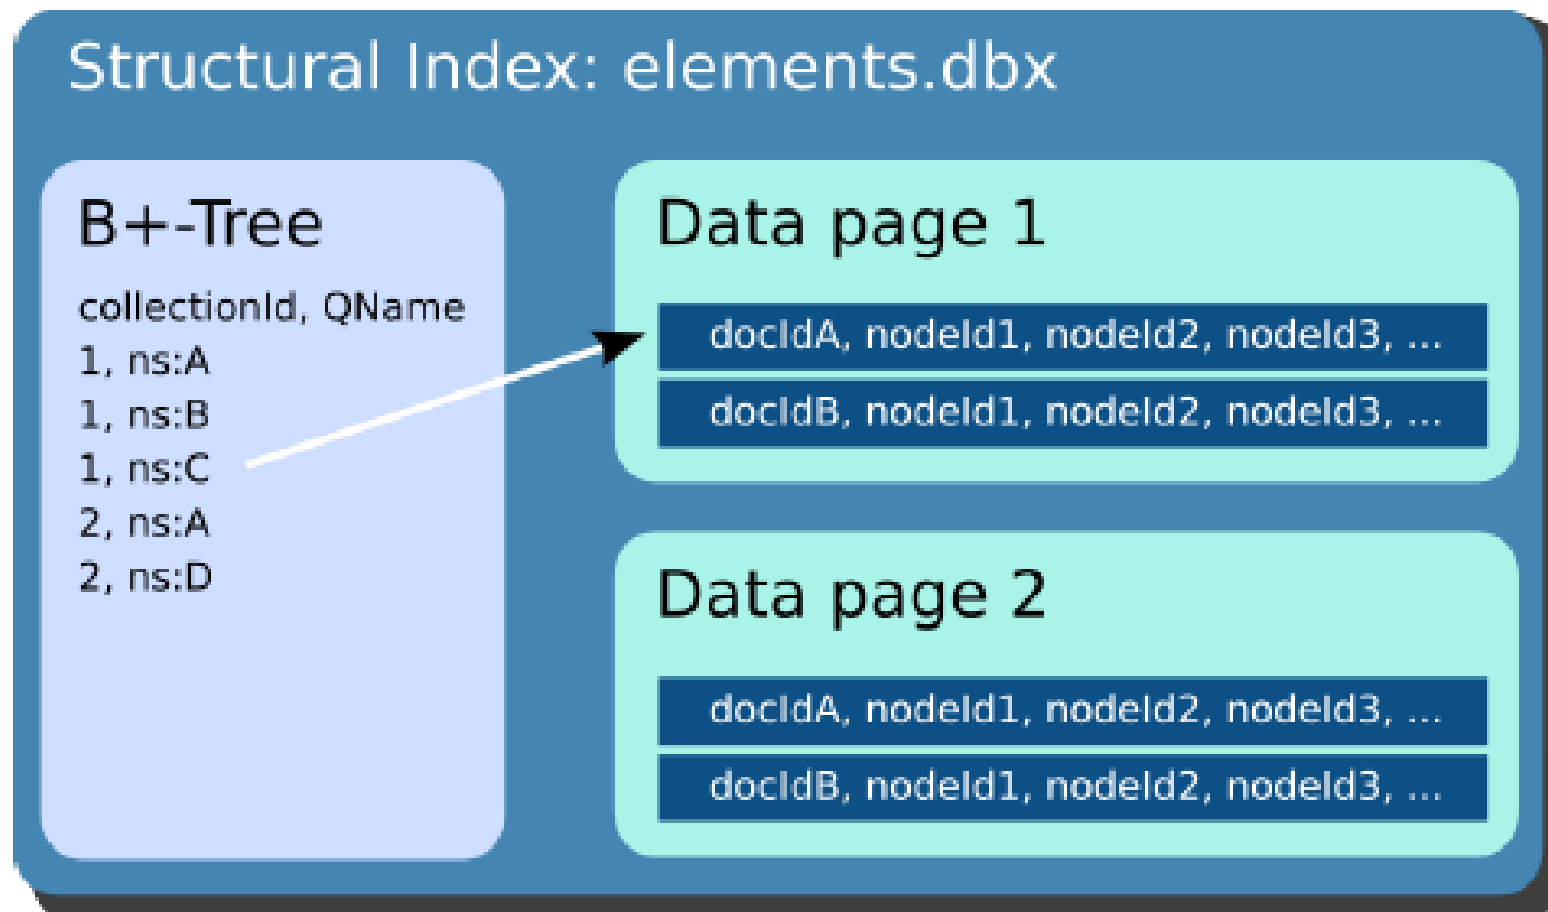
dom.dbx

Every node in eXist is identified by a tuple **<docId, nodeId>**




elements.dbx: Structural Index

- Indexed by **<collection-id, name-id>** pairs
- Value is ordered list of **<document-id, node-id>** pairs, which are the nodes whose name matches the key entry



elements.dbx: Structural Index



Find all the documents in collection c that have an element “phone”.

Query engine

- eXist prefers **XPath** predicate expressions over an **equivalent FLWOR construct using a "where" clause!**

```
for $i in //entry
where $i/@type = 'subject' or $i/@type = 'definition' or
      $i/@type = 'title'
return $i
```

```
//entry[@type = ('subject', 'definition', 'title')]
```

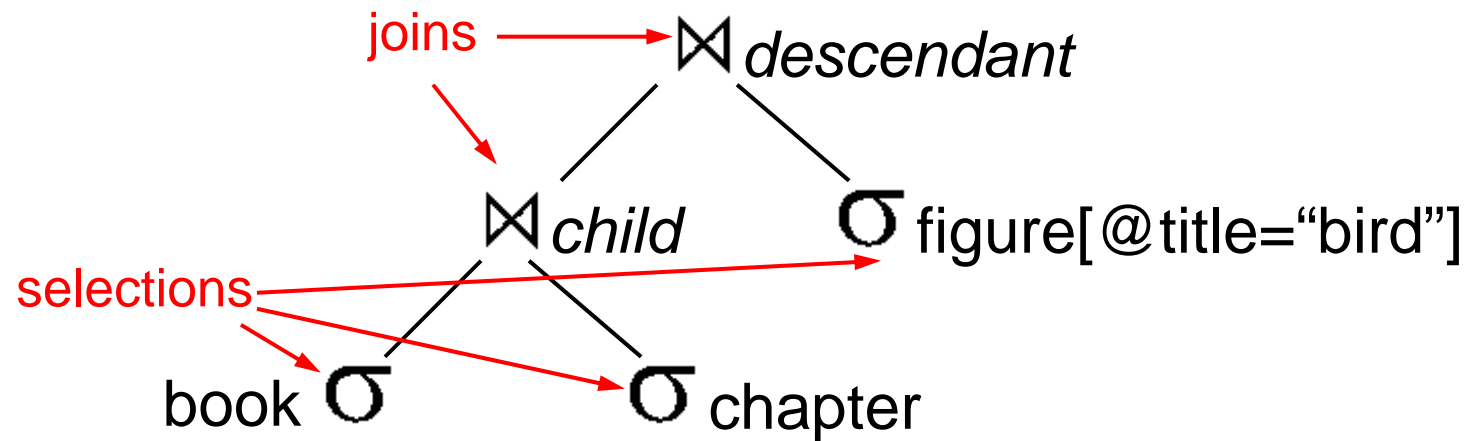
- Based on these provided index types, the eXist XQuery engine relies on **path join algorithms** for efficient computation of node relationships instead of traditional tree traversals.

Limitations

- Currently, most optimizations are implicit
- “Query plan” is hard-coded into the query engine
- Hard to maintain/debug/profile
- eXist needs a better, query-rewriting optimization engine!


Evaluating path expressions

- `//book/chapter//figure[@title="bird"]`



Join algorithms

//book/chapter



- **Nested loop:**

1. Find “book” elements by enumerating all elements
2. For each “book” element, enumerate all children and output the ones named “chapter”

- **“Path join”:** efficient join operations utilizing indexes

1. $A = \{\text{IDs of all elements named “book”}\}$, obtained through the **elements.dbx**
2. $B = \{\text{IDs of all elements named “chapter”}\}$
3. *child axis join*: $\text{Result} = \{x \mid x \in B, \text{parent}(x) \in A\}$

Annexe

Example of how indexes used

/play//speech[speaker = 'Hamlet']


/play//speech



speech[speaker]



speaker = 'Hamlet'

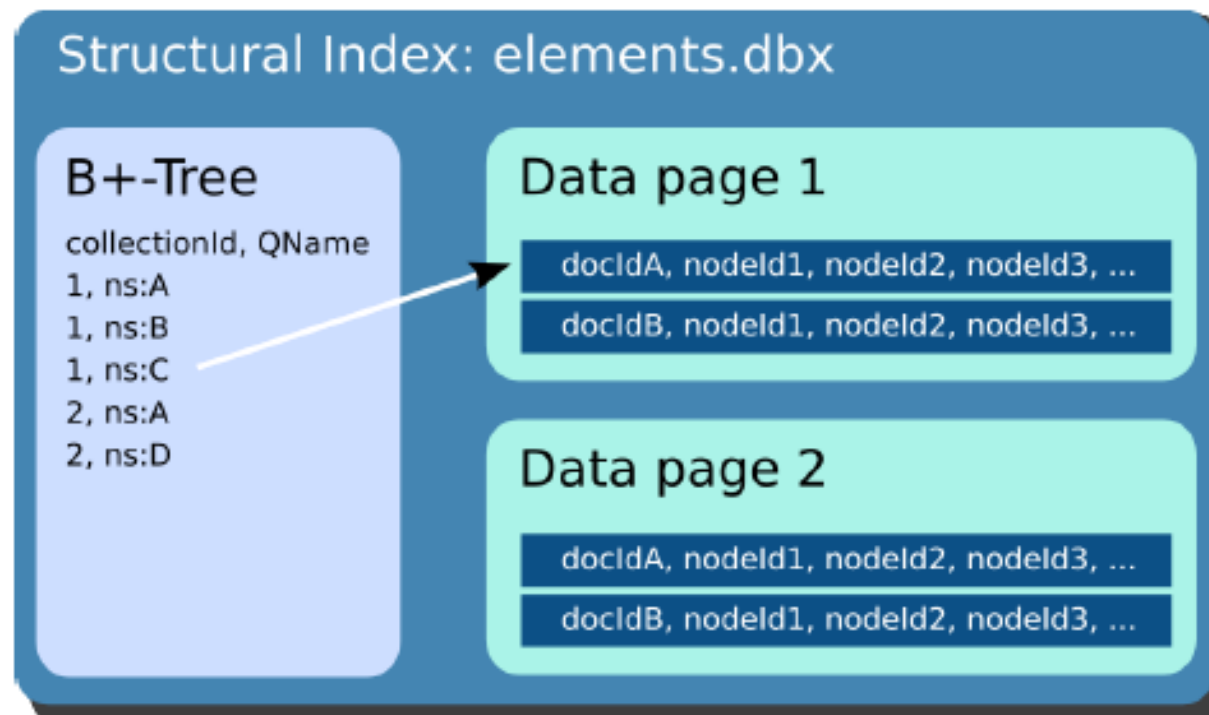


Annexe Example of how indexes used: First step

```
/play//speech[speaker = 'Hamlet']
```

```
/play//speech
```

- find **play** element for all documents using **elements.dbx** → set₁ <doc_id, node_id>
- find **speech** elements using **elements.dbx** → set₂ <doc_id, node_id>



Annexe Example of how indexes used: First step

/play//speech[speaker = 'Hamlet']

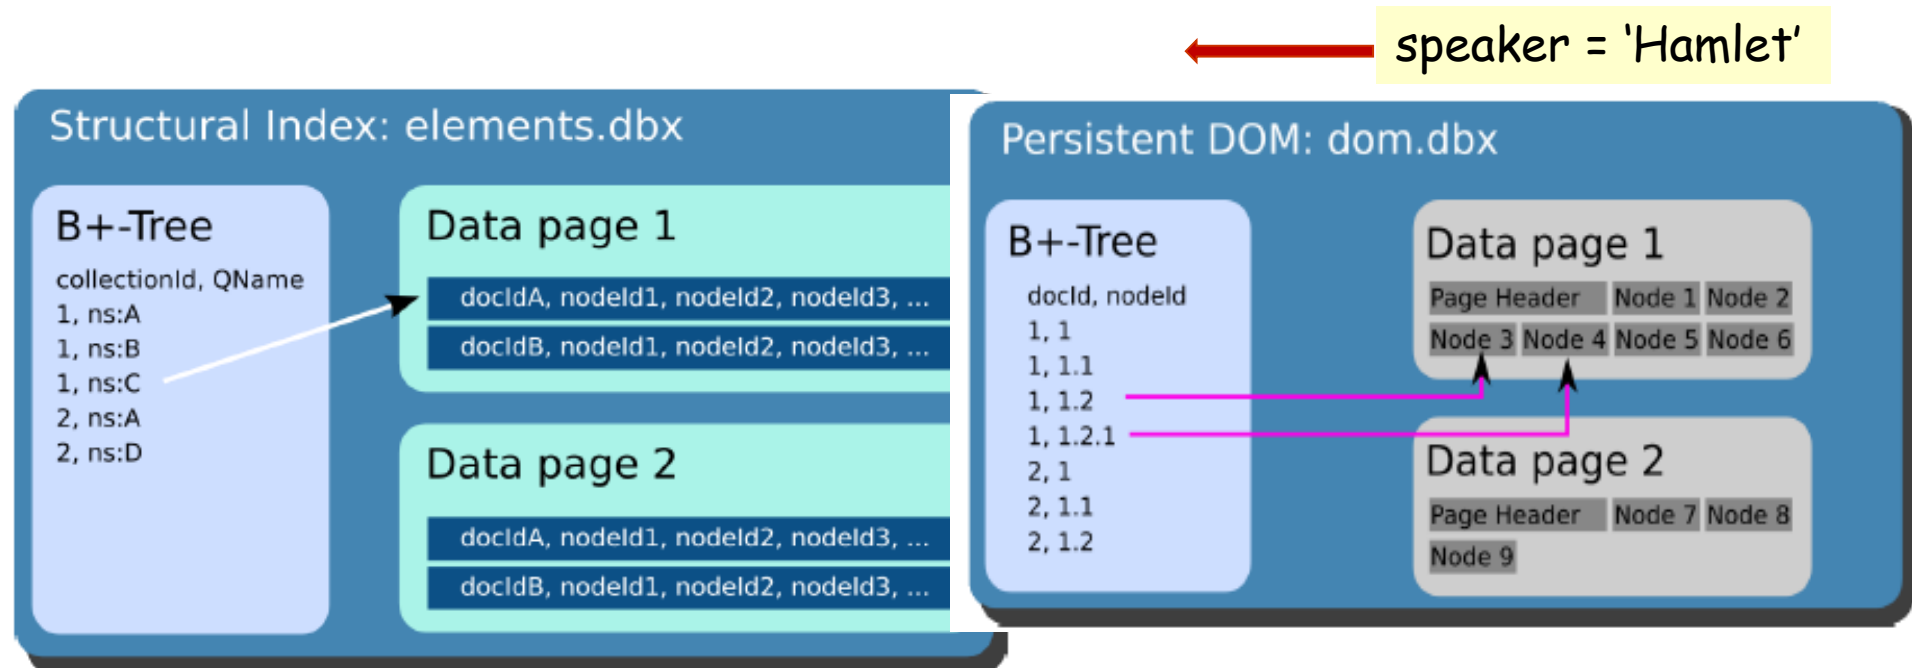
/play//speech → speech[speaker]

- find **play** element for all documents using **elements.dbx** → set₁ <doc_id, node_id>
- find **speech** elements using **elements.dbx** → set₂ <doc_id, node_id>
- Ancestor-descendant join algorithm is applied to set₁ and set₂ (<document-id, node-id>): keep the **speech** which are descendant of **play** → we also get the ancestor node set for the expression **speech[speaker]**

```
<div6 exist:id="1.46.9.7.8" xml:id="JG10229">
  <head exist:id="1.46.9.7.8.7">
    <name exist:id="1.46.9.7.8.7.3" type="Dramenfigur">FA
    <stage exist:id="1.46.9.7.8.7.4">unruhig</stage>
    <stage exist:id="1.46.9.7.8.7.6"> auf seinem Sessel a
  </head>
  <sp exist:id="1.46.9.7.8.8" xml:id="JG10230">
```


Annexe Example of how indexes used: Second step

/play//speech[speaker = 'Hamlet']



- find **speaker** element for all documents using **elements.dbx** and use **dom.dbx** to access the actual DOM nodes in the XML store and keep $\rightarrow \text{set}_4 \langle \text{doc_id}, \text{node_id} \rangle$ the speaker nodes having value "Hamlet"

Annexe

Example of how indexes used: Last step

```
/play//speech[speaker = 'Hamlet']
```

`/play//speech` \longrightarrow `speech[speaker]` \longleftarrow `speaker = 'Hamlet'`

- find **play** element for all documents \rightarrow $\text{set}_1 \langle \text{doc_id}, \text{node_id} \rangle$
- find **speech** elements \rightarrow $\text{set}_2 \langle \text{doc_id}, \text{node_id} \rangle$
- Ancestor-descendant join algorithm is applied to keep the **speech** which are descendant of **play** \rightarrow $\text{set}_3 \langle \text{doc_id}, \text{node_id} \rangle$
- find **speaker** element having value "Hamlet" \rightarrow $\text{set}_4 \langle \text{doc_id}, \text{node_id} \rangle$
- Finally, the resulting set **set₄** is joined with the context node **set₃** by applying an ancestor-descendant path join and output the resulting nodes **speech**.