

Module: XML et les bases de données

Les SGBD natifs XML (XDB)

Houda Chabbi Drissi

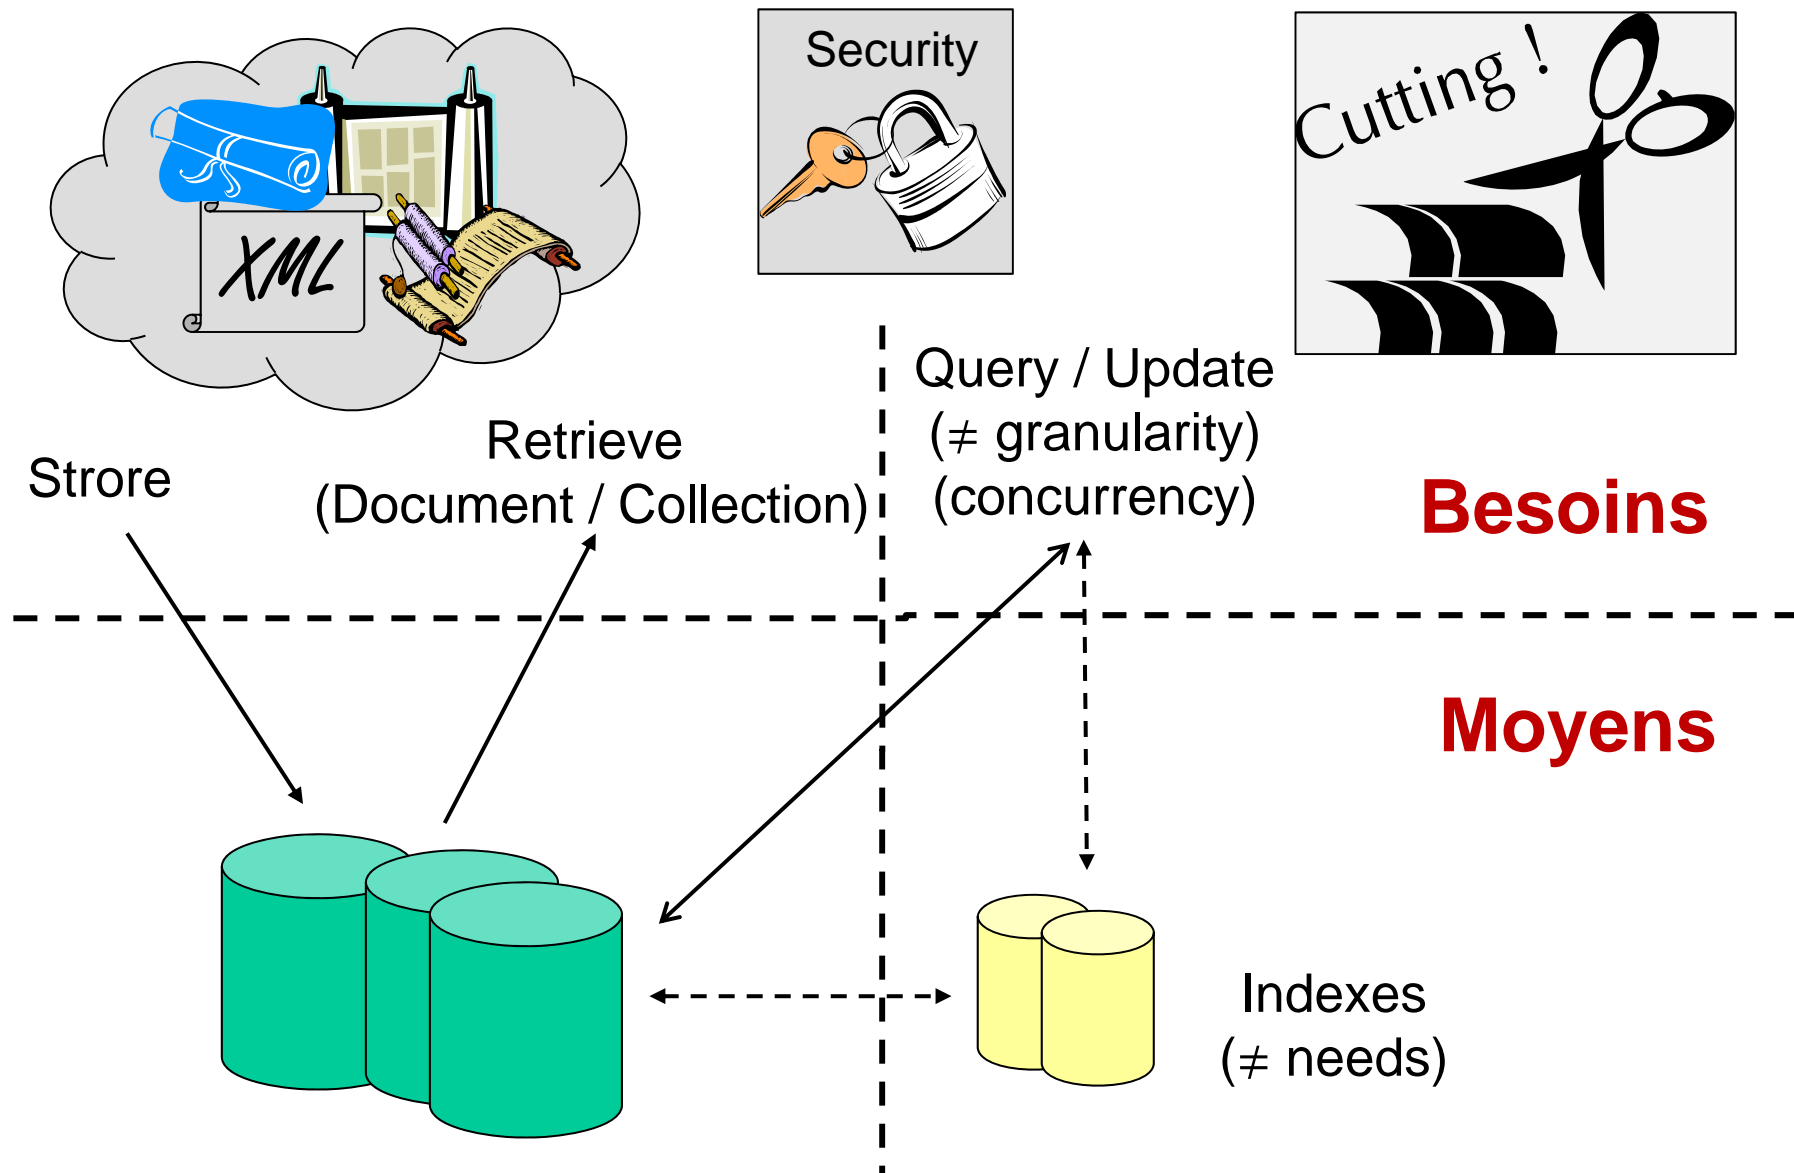
houda.chabbi@hefr.ch

Plan

► Motivation – Needs – Means.

- NXD Concepts
- Indexing approaches
- Conclusion

Motivation



Users' database needs

- Store/Retrieve - global/parts of XML
- Query/Update - global/parts of XML
- Independency: Logical model versus Physical model
- Optimization
- ACID Transactions
- Recovery
- Security

Means

- Choose logical/physical models for storage/retrieval's mechanism.
- Query/Update languages.
- Optimization: *indexing*, caching.
- Concurrency handling : ACID.
- Recovery approaches.

Databases' solutions

Plan

- Motivation – Needs – Means.
- ▶ NXD Concepts
 - Indexing approaches
 - Conclusion

Native XML Data Stores

- Provides a **logical model** for storing and retrieving XML documents efficiently
- XML interface for accessing XML data (e.g. **XPath, XQuery**)
- Advantages
 - Scalability
 - Data-access speed
 - Reliability

Native XML Database

- Definition (XML:DB Initiative – www.xmldb-org.sourceforge.net)
 - XML document is the fundamental unit of logical storage
 - Uses a logical model for the XML document itself
 - Requires a particular underlying physical storage model

Abstract View of Components

- Native XML Interface
 - Most popular choices are XQuery and Xpath

- Storage Manager
 - Performs any necessary transformations to store/retrieve XML data

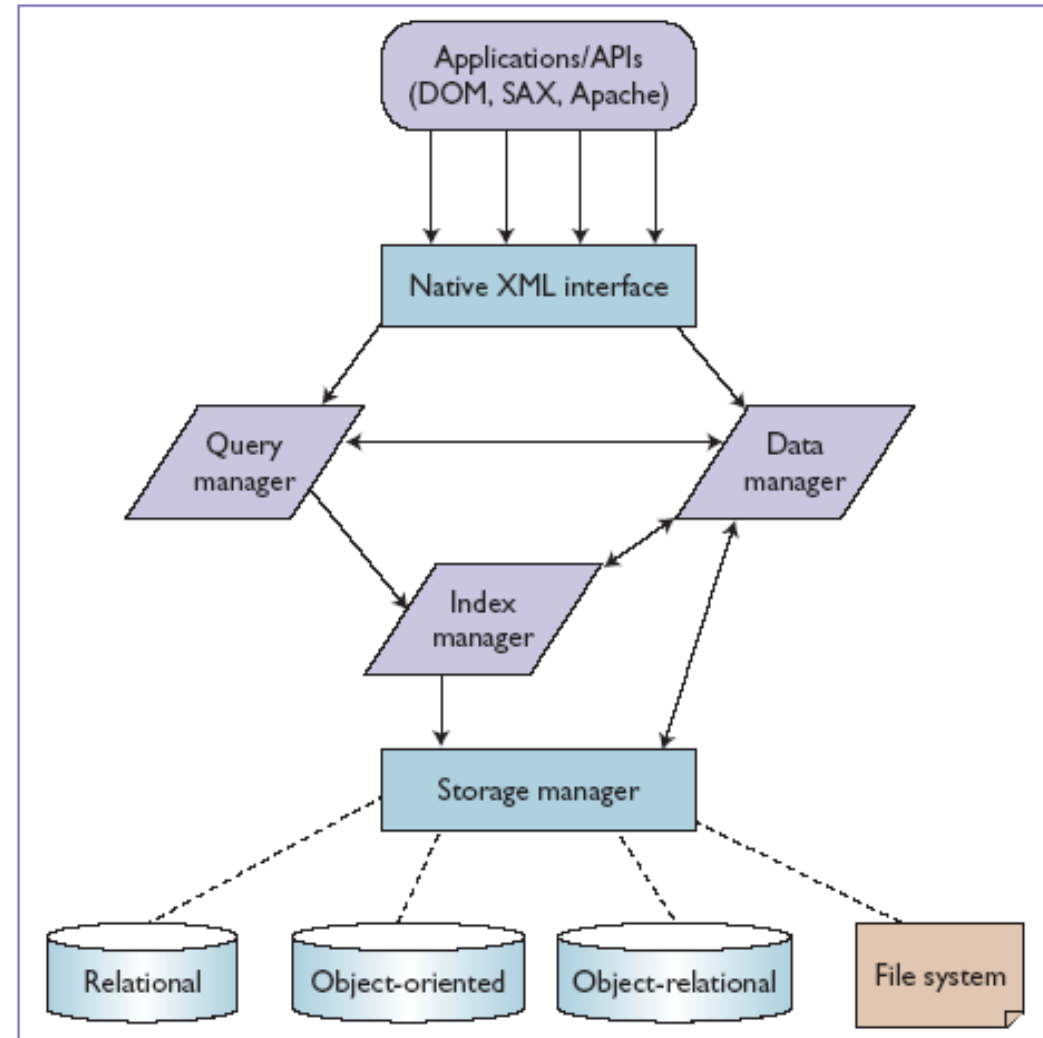
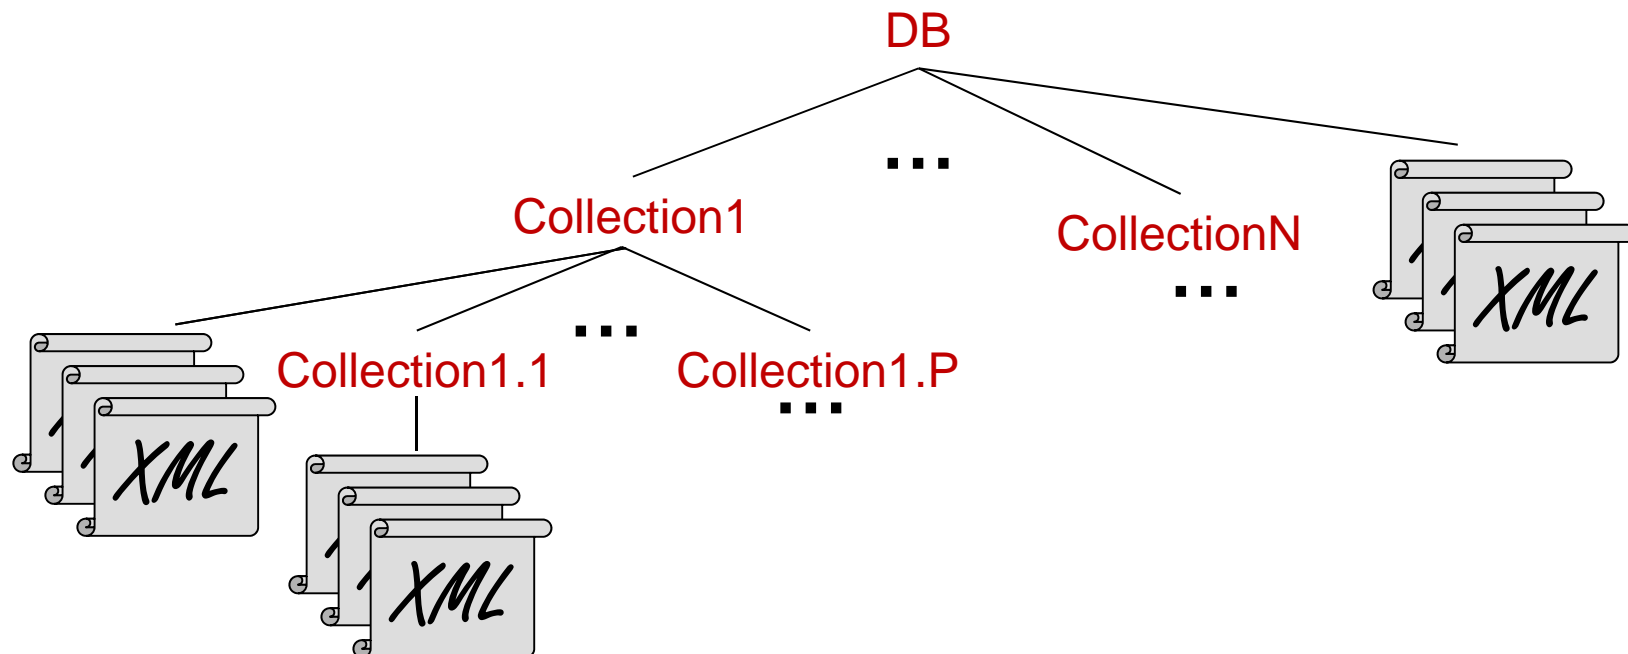


Figure 1. Abstract view of a native XML data management framework. The store includes two key components: a native XML interface and a storage manager.

Vocabulary: XML Collections

- **XML Collection** is a collection of XML documents
- “A row is to a table as an XML document is to an **XML collection**.”



Collections

NXDs manage collections of documents, allowing to query and manipulate those documents as a set.
→ similar to the relational concept of a table.

NXDs diverge from the table concept in that not all native XML databases require a schema to be associated with a collection: any XML document can be store in the collection, regardless of schema.

→ can construct queries across all documents in the collection.

Schema-independent NXDs

☺ Having **schema-independent document collections** gives the database a lot of flexibility and makes application development easier.

☹ Unfortunately, the **risk of low data integrity** is high.

- If strong schema structure is needed, then:
- either make sure you use a NXD that supports schemas
 - or find another way to store your XML data.

Queries / Updates: standard languages

- **Xquery** is the NXD query language
- **Xupdate** is the NXD update language
- **Xquery Full Text** fusion between IR and databases concepts.

Plan

- Motivation – Needs – Means.
- NXD Concepts
- ▶ Indexing approaches
- Conclusion

Indexing characteristics

- Index Type:
 - ▶ Structure + implementation example
 - value,
 - full-text.
- Index Scope:
 - document, collection.
- Index Target:
 - document, node.
- Index Control:
 - automatic, voluntary, required.

Index type: structural indexes or path indexes

Structural indexes materialize results of path expressions. Useful for navigational portions of queries.

Used to answer to *PATH*, such as *"/a/b/c"*.

Implementation: *A(k)* indexes

"/a/b/c"

<a>	1
foo	1.1
foo	1.2(.1)
	1.3
<c>foo</c>	1.3.1(.1)
	1.3.2
bar	1.3.2.1
<a>	1.3.2.2
bar	1.3.2.2.1
	1.3.2.2.2
	
	
foo	1.4(.1)
<c>bar</c>	1.5(.1)
	

Figure 1: Example XML document. Dewey order encoding of elements shown on the right.

1. /a
2. /a/b
3. /a/b/c
4. /a/b/b
5. /a/b/b/a
6. /a/b/b/a/b
7. /a/c

Figure 2: Enumeration of unique paths seen in XML in Figure 1. This is the set of paths which would be indexed by a path index.

<a>	1
foo	1.1
foo	1.2(.1)
	1.3
<c>foo</c>	1.3.1(.1)
	1.3.2
bar	1.3.2.1
<a>	1.3.2.2
bar	1.3.2.2.1
	1.3.2.2.2
	
	
foo	1.4(.1)
<c>bar</c>	1.5(.1)
	

Figure 1: Example XML document. Dewey order encoding of elements shown on the right.

" /a//c " ?

Les c qui ont pour ancêtre a?

a ancêtre de c donc position a avant c

→ Reste: path 3 et path 7

1. /a
2. /a/b
3. /a/b/c
4. /a/b/b
5. /a/b/b/a
6. /a/b/b/a/b
7. /a/c

a: 1,1 2,1 3,1 4,1 5,1 5,4 6,1 6,4 7,1
b: 2,2 3,2 4,2 4,3 5,2 5,3 6,2 6,3 6,5
c: 3,3 7,2

Figure 4: c dans path 3 à la profondeur de 3 et path 7 à la profondeur de 2

Figure 2: Enumeration of unique paths seen in XML in Figure 1. This is the set of paths which would be indexed by a path index.

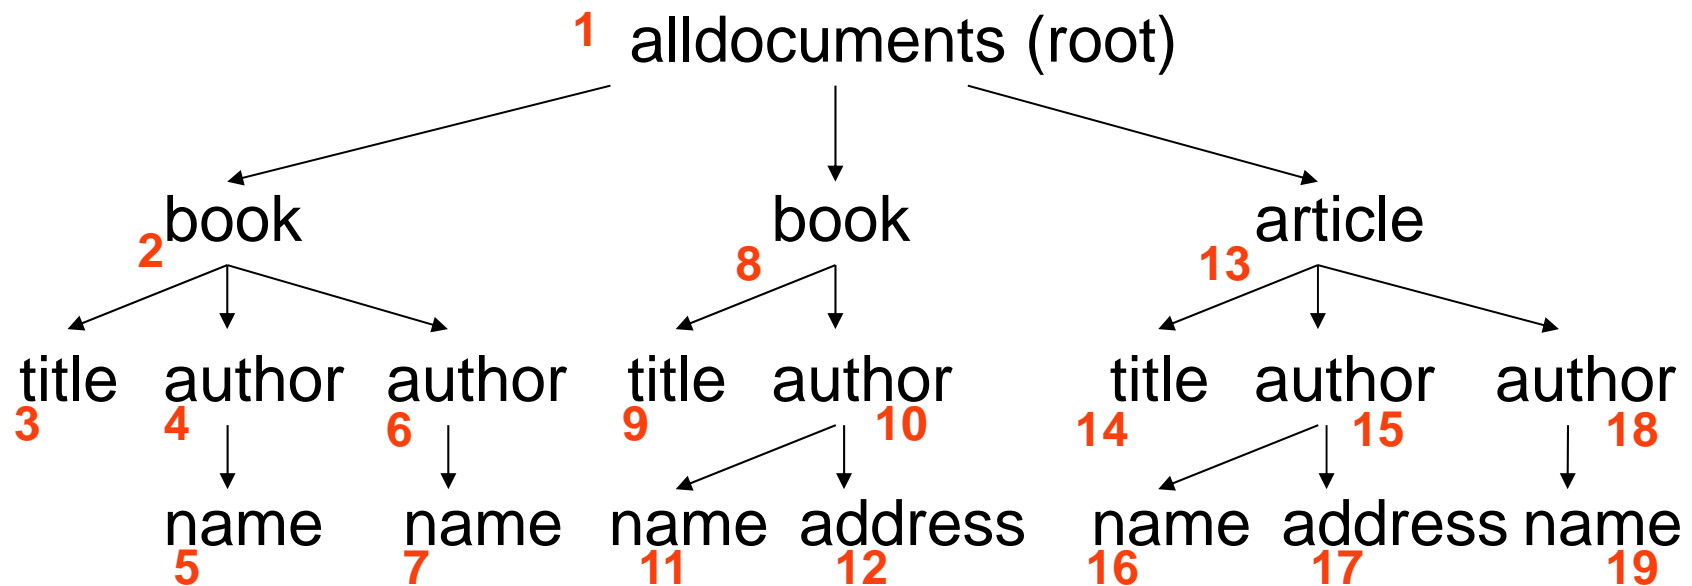
- Two nodes u, v are called **bisimilar** if:
 - They have the same label.
 - If u' is the parent of u , then there is a parent v' of v , such that u', v' are also bisimilar, and vice versa.

⇒ two nodes are **bisimilar** *if for every finite path starting at one node there is a path with identical sequence of labels starting from the other.*

Définition récursive. En fait, les deux nœuds sont atteignables par un même chemin depuis la racine.

Annexe

XML example: data graph

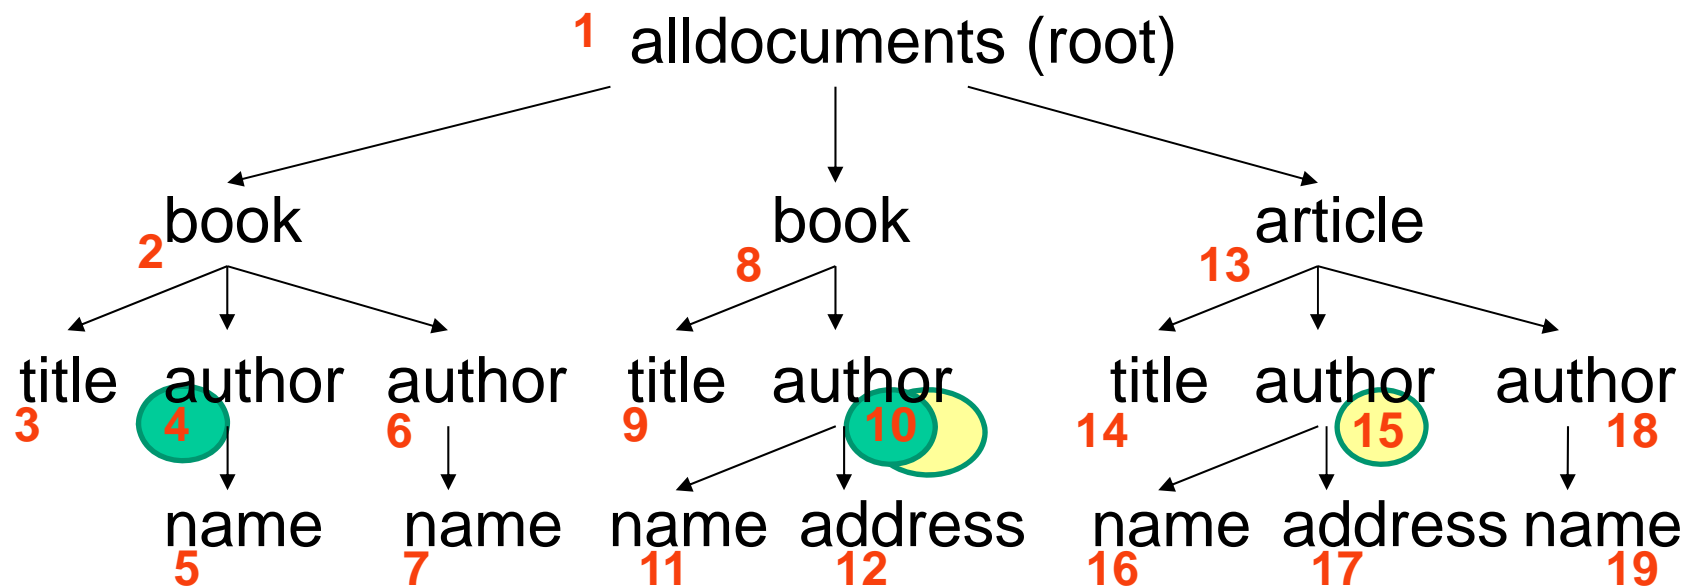


Source: http://i.cs.hku.hk/~nikos/courses/CSIS7101/intro_slides.pdf

Annexe

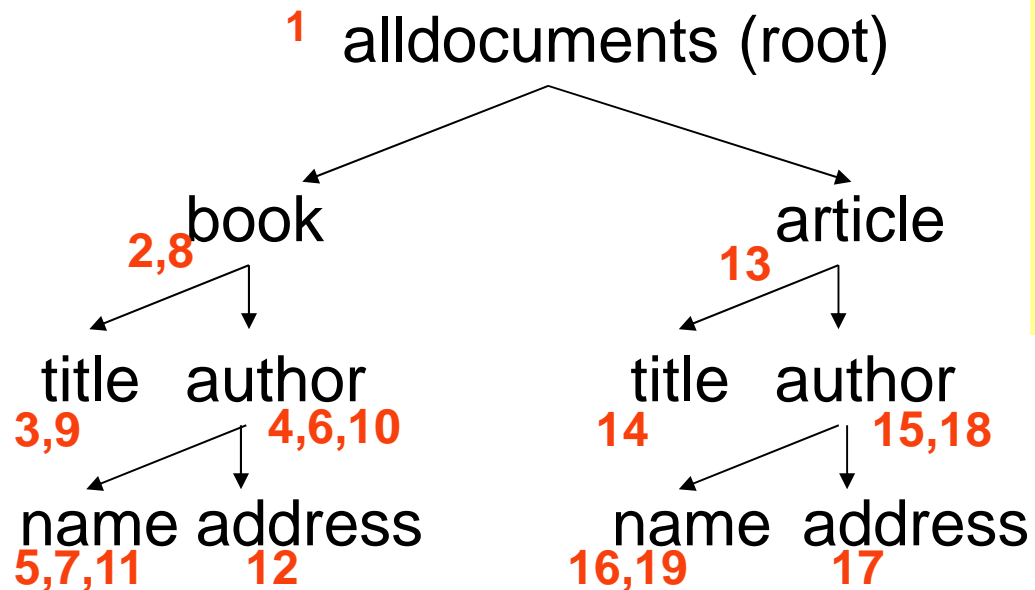
Bisimilarity examples

- Nodes 4 and 10 are bisimilar
- Nodes 10 and 15 are not bisimilar



Annexe

1-index graph



/alldocuments/book/title
/alldocuments/book/author
/alldocuments/book/author/name
/alldocuments/book/author/address
/alldocuments/article/title
/alldocuments/articleauthor
/alldocuments/article/author/name
/alldocuments/article/author/address

- Index summarizes path information
- Each entry: list of pointers to data nodes

- The **1-index** maintains information about **all paths** in the original graph.
- This makes **the index very large** (with size comparable to the data size) \Rightarrow quite expensive to evaluate queries using this index.
- To address this problem an **A(k)-index** is proposed which **indexes exactly only paths up to length k**.

- Two nodes u, v are 0-bisimilar, if they have the same label.

- Two nodes u, v are k -bisimilar,
 - if they have the same label
 - and for every parent u' of u , there is a parent v' of v , such that u' and v' are $(k-1)$ -bisimilar, and vice versa.

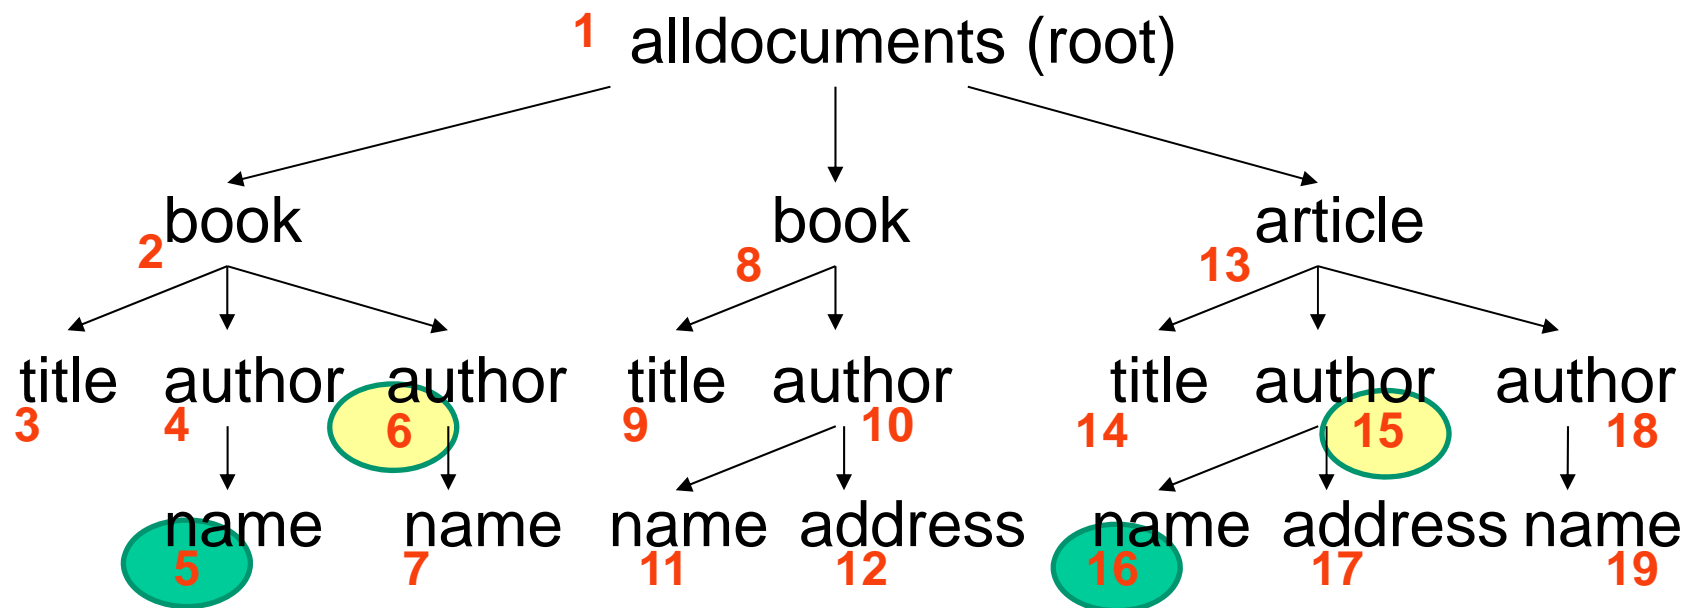
Properties

- If **a** and **b** are **bisimilar**
 - set of incoming paths into them is same (from **root**)
- If **a** and **b** are **0-bisimilar**
 - They have the same label
- If **a** and **b** are **k-bisimilar**
 - set of incoming paths of length $\leq k$ are same

Annexe


k-bisimilarity Example

- Nodes 5 and 16 are 1-bisimilar
- Nodes 6 and 15 are not 1-bisimilar



- The $A(k)$ -index stores exactly all paths of length k , or else: all k -bisimilar nodes in the data graph are stored in the same node in the index graph.
- This means that all incoming paths up to length k are encoded in the index.

XML example: data graph + 0-bisimilar



Annexe

XML example: data graph + 1-bisimilar

alldocument/book

book/author

article/author

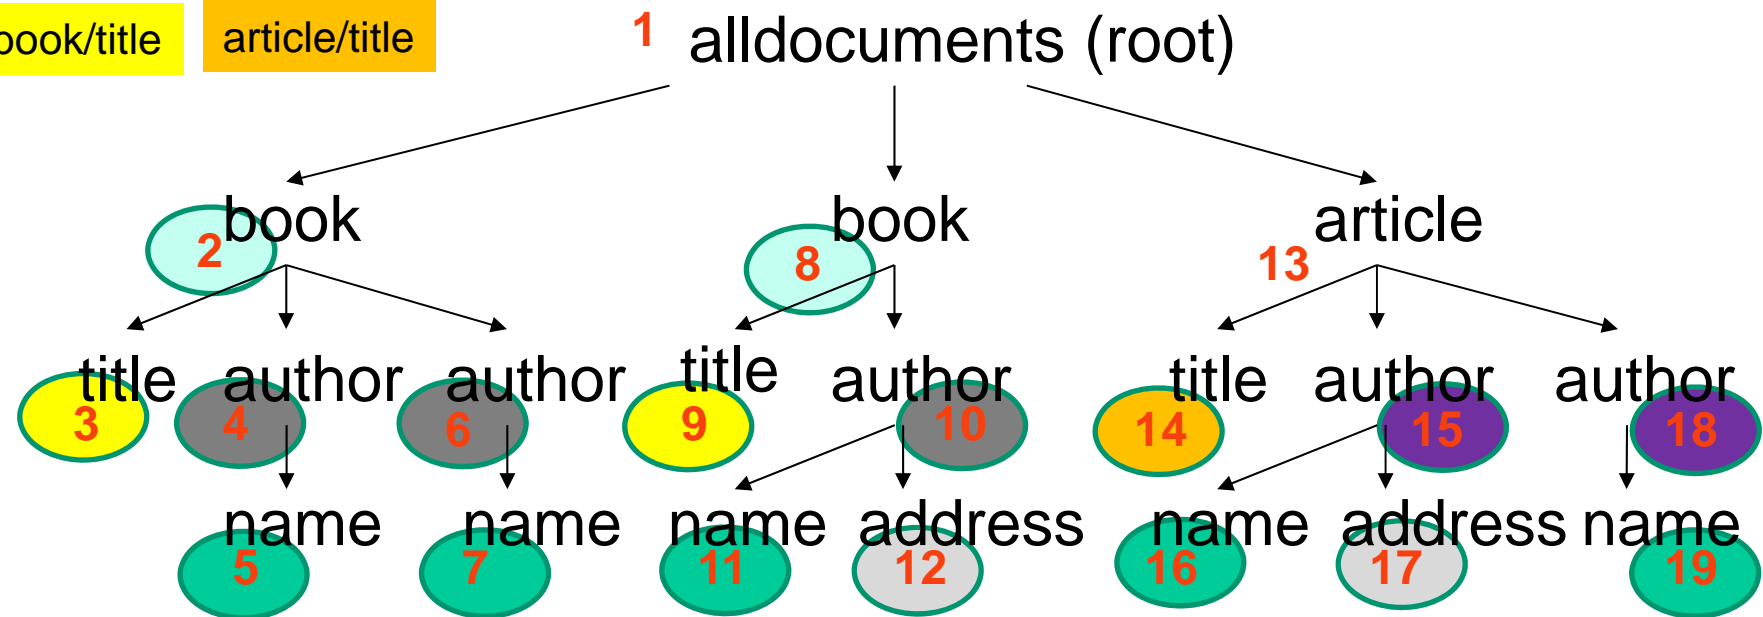
author/name

author/address

alldocuments/article

book/title

article/title



Source: http://i.cs.hku.hk/~nikos/courses/CSIS7101/intro_slides.pdf


Annexe

A(k)-index examples

book
author
name
address
article
title

A(0)-index


1 alldocuments (root)



alldocument/book
book/author
author/name
author/address
alldocuments/article
book/title
article/title

A(1)-index

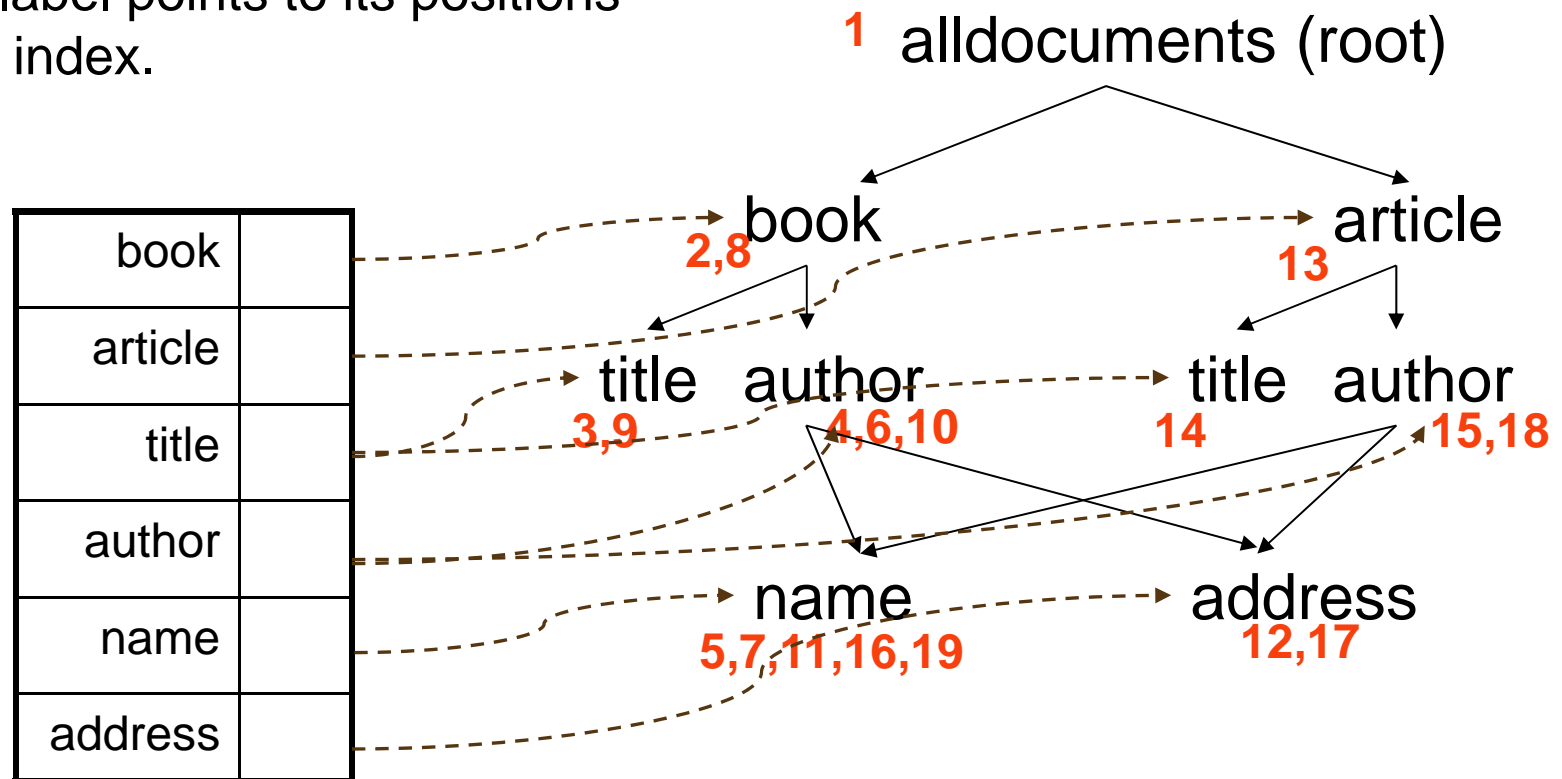
1 alldocuments (root)



Annexe

Using the A(k) index to search

A *Label Map* is constructed together with the index, where each label points to its positions in the index.

A(1)-index

Source: http://i.cs.hku.hk/~nikos/courses/CSIS7101/intro_slides.pdf


Assume that a path query q of length $\leq k$ is applied.
The $A(k)$ index can answer the query as follows:

1. First the last label of q is found and the label map is used to **find its positions in the index**.
2. Then the index **is traversed backwards** to complete the answer.

Evaluation of path queries (example)

- Query: **book/title**

A(1)-index



Problems of path indexes

- They are appropriate only for **simple path queries** up to a **certain length**.
- Therefore if a query has branches or regular path expressions the index cannot provide exact answers, but the actual data have to be accessed.
- Also **these indexes have high storage and update cost**.

Indexing characteristics

- Index Type:
 - Structure + implementation example
 - ▶ value,
 - full-text.
- Index Scope:
 - document, collection.
- Index Target:
 - document, node.
- Index Control:
 - automatic, voluntary, required.

Index type: value indexes

Value indexes: may be typed \Rightarrow Useful for comparisons (=, <, etc.).

Track all values for specific **elements** or **attributes**: (value, label)

- index atomic values; e.g., data(//emp/salary)
- use **B+ trees**

A value index on the element "color" would have an index entry for every separate instance of color and would be useful for a query such as: **//color[.='green']**.

```

<a>                                1
  foo                             1.1
  <b>foo</b>                       1.2(.1)
  <b>                               1.3
    <c>foo</c>                     1.3.1(.1)
    <b>                             1.3.2
      bar                         1.3.2.1
      <a>                           1.3.2.2
        bar                     1.3.2.2.1
        <b>                       1.3.2.2.2
      </a>
    </b>
  </b>
  <b>foo</b>                       1.4(.1)
  <c>bar</c>                       1.5(.1)
</a>

```

Figure 1: Example XML document. Dewey order encoding of elements shown on the right.

→ foo dans uniquement Path 3

```

"foo": 1 1.1
       2 1.2.1
       3 1.3.1.1
       2 1.4.1
"bar": 4 1.3.2.1
       5 1.3.2.2.1
       7 1.5.1

```

Figure 3: Value index for example XML from Figure 1. Storing path ID and path instance Dewey order. Document ID and position omitted.

" /a//c/'foo' "

1. /a
2. /a/b
3. /a/b/c
4. /a/b/b
5. /a/b/b/a
6. /a/b/b/a/b
7. /a/c

```

a: 1,1 2,1 3,1 4,1 5,1 5,4 6,1 6,4 7,1
b: 2,2 3,2 4,2 4,3 5,2 5,3 6,2 6,3 6,5
c: 3,3 7,2

```

Figure 4: Inverted lists for paths seen in example XML. Storing path ID and position within path.

→ /a//c/ Path possibles 3 ou 7

Figure 2: Enumeration of unique paths seen in XML in Figure 1. This is the set of paths which would be indexed by a path index.

Indexing characteristics

- Index Type:
 - Structure + implementation example
 - value,
 - ▶ full-text.
- Index Scope:
 - document, collection.
- Index Target:
 - document, node.
- Index Control:
 - automatic, voluntary, required.

Index type: Full-text indexing

⇒ Related to IR domain

Full Text indexes:

- Keyword search, inverted files
- IR world, text extenders

Some XML database products implement what they call full-text indexing, which minimally is a word index over a document.

Exemple

Doc 1:

```
<invoice>
  <buyer>
    <name>ABC Corp</name>
    <address>1 Industrial
Way</address>
  </buyer>
  <seller>
    <name>Acme Inc</name>
    <address>2 Acme Rd.</address>
  </seller>
  <item count=3>saw</item>
  <item count=2>drill</item>
</invoice>
```

Doc 2:

```
<invoice>
  <buyer>
    <name>Oracle Inc</name>
    <phone>555-1212</phone>
  </buyer>
  <seller>
    <name>IBM Corp</name>
  </seller>
  <item>
    <count>4</count>
    <name>nail</name>
  </item>
</invoice>
```

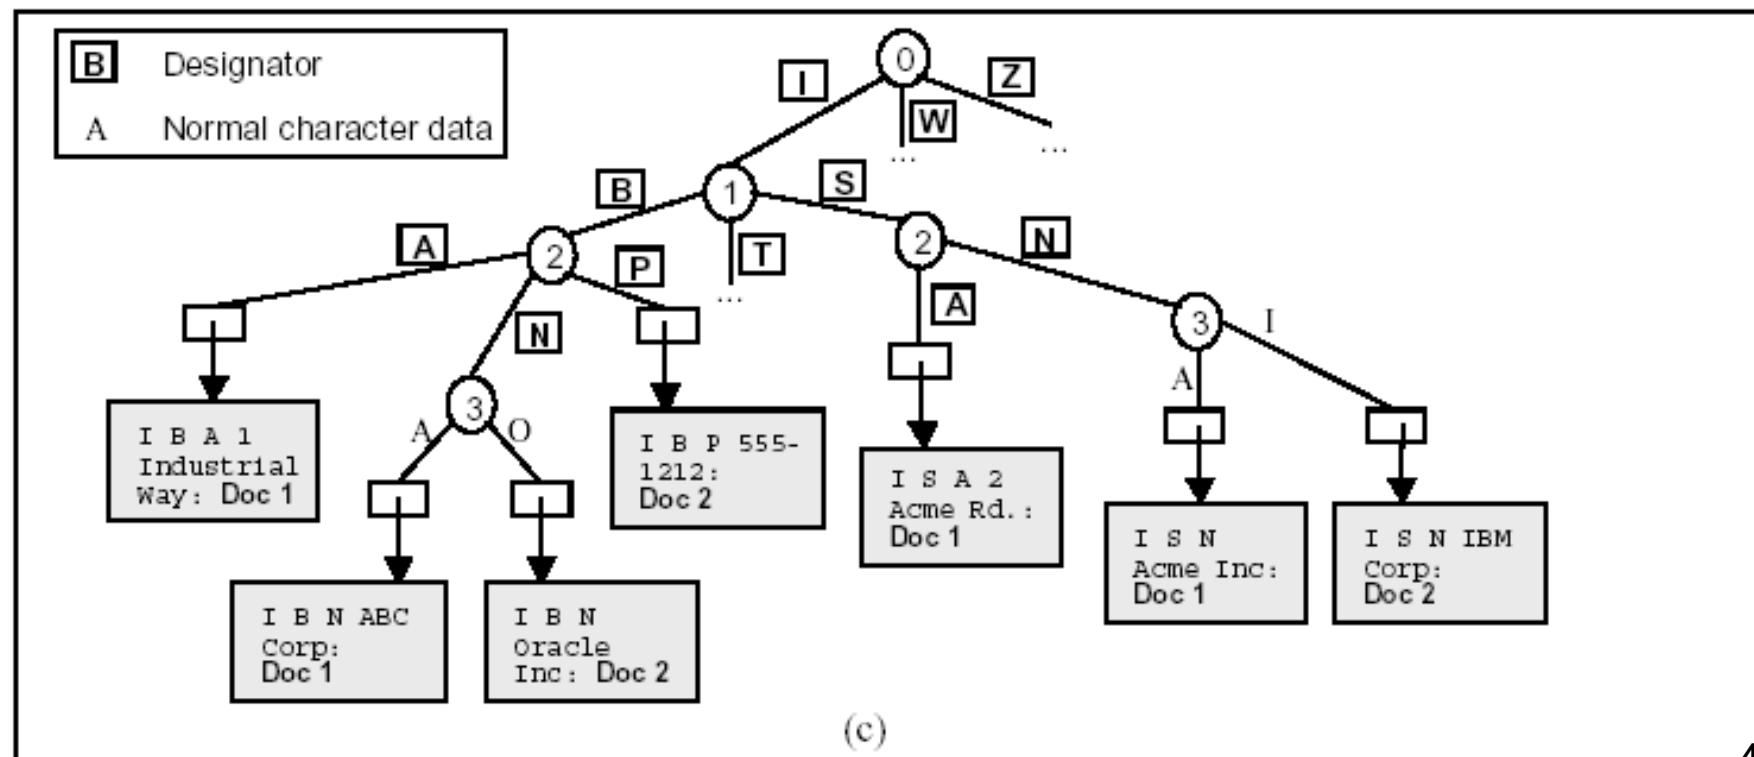

Annexe

Patricia Trie

(a) <invoice> = I
 <buyer> = B
 <name> = N
 <address> = A
 <seller> = S
 <item> = T
 <phone> = P
 <count> = C
 count (attribute) = C'

(b)

Document 1	Document 2
I B N ABC Corp	I B N Oracle Inc
I B A 1 Industrial Way	I B P 555-1212
I S N Acme Inc	I S N IBM Corp
I S A 2 Acme Rd.	I T C 4
I T drill	I T N nail
I T C' 2	
I T saw	
I T C' 3	



Indexing characteristics

- Index Type:
 - structure, value, full-text.
- ▶ Index Scope:
 - document, collection.
- Index Target:
 - document, node.
- Index Control:
 - automatic, voluntary, required.

Index scope

Most native XML databases store **documents** in a **collection**. The scope of a given index could be:

- **collection-wide**
- **or it could be restricted to a single document.**

A native XML database system can choose the index scope it implements. Queries against a collection can return documents or sets of nodes within documents. In order to support efficient restriction of a query to a manageable set of documents, the system must support indexes at the collection scope.

Indexing characteristics

- Index Type:
 - structure, value, full-text.
- Index Scope:
 - document, collection.
- ▶ Index Target:
 - document, node.
- Index Control:
 - automatic, voluntary, required.

Index target

Related to scope is the target or the object referenced by an index entry. It can be:

- a document
- or an object within a document.

An index is capable of pointing down to the addressable unit in the system, but such granularity is not always necessary and can be expensive.

Although this is possible, it is the case that most database systems with fine-grained document storage reference directly to nodes in indexes rather than to the containing documents.

Indexing characteristics

- Index Type:
 - structure, value, full-text.
- Index Scope:
 - document, collection.
- Index Target:
 - document, node.
- ▶ Index Control:
 - automatic, voluntary, required.

- **Voluntary indexes** are specified explicitly by an interface to the system. These indexes allow for some experimentation to find the minimal useful set of indexes.
- **Some systems have automatic indexes**, where a well-defined set of indexes always is created, except for those that are disabled explicitly, by way of configuration or interface. The system also may have required indexes, which cannot be disabled because they are necessary for proper functioning of the system.

Plan

- Motivation – Needs – Means.
- NXD Concepts
- Indexing approaches
- ▶ Conclusion

Conclusion

- Logical model: **collection-documents**
- Physical model: depends on the products
- Query/Update: **Xquery (Xpath)**
- Optimization: **Indexes and execution plan of queries**
- ACID Transactions: like for traditional DBMS (BaseX)
- Security: **users accounts**
- Recovery: like for traditional DBMS